

Evaluating Time-Series Foundation Models

1 Question

This project evaluates how foundation-model forecasts on the official TIME test tasks change with target representation, usable covariates, input normalization, and available context length. Each comparison holds the TIME dataset windows and metric implementation fixed unless its protocol explicitly declares another factor.

2 Current migrated comparisons

The initial executable surface contains a five-model foundation benchmark and a Chronos-2 comparison between native multivariate targets, independent univariate targets, and other target histories supplied as past-only covariates. Seasonal Naive is deterministic and provides the taskwise MASE normalizer. It also fixes the evaluation grid to cells with finite target support, finite Seasonal Naive predictions on that support, and finite Seasonal Naive MASE. Every learned model uses this same grid; a non-finite forecast on expected support invalidates its task rather than changing the aggregation support. Inference timing excludes loading, dataset construction, metric computation, and artifact saving.

3 Evidence boundary

The repository makes no result claim until its migrated experiments are rerun and their project-owned manifests and aggregates are inspected. Covariate, normalization, and context-size extensions remain planned until their exact grids are implemented.